



Digital Surveillance Systems and Application

Coding manual


Jison G.S. Hsu

Artificial Vision Laboratory

Taiwan Tech

Example 7.1 : Face Recognition

- Please use the dlib models on Colab to recognize the faces in image and show the result with detection boxes as shown below



Sol 7.1 : Face Recognition (1/10) – Face Verification

- Face verification is a 1: 1 comparison.
- The identity verification mode is essentially a process by which the computer quickly compares the current face with the portrait database and determines whether it matches.

```
!pip install face_recognition
!pip install pillow
!pip install numpy
```

```
#face detection part
import face_recognition
image = face_recognition.load_image_file('./two_people.jpg') #load image
face_locations = face_recognition.face_locations(image) #detect face


#Array of coords of each face
print(face_locations)
print(f'There are {len(face_locations)} people in this image')
```

```
#face recognition part
import face_recognition
obama_1=face_recognition.load_image_file('./obama.jpg')
#feature extractor
obama_1_encoding = face_recognition.face_encodings(obama_1)[0]

unknown_im=face_recognition.load_image_file('./biden.jpg')
unknown_im_encoding = face_recognition.face_encodings(unknown_im)[0]

# feature compare
results = face_recognition.compare_faces([obama_1_encoding],unknown_im_encoding)
if results[0]:
    print('This is obama')
else:
    print('This is not obama')
```

Test Pair



0.62

similarity

0.3



True

→ Threshold: 0.5

False

Sol 7.1 : Face Recognition (2/10)

```
# face detection and show the part of face
from PIL import Image
import face_recognition
image = face_recognition.load_image_file('./two_people.jpg')
face_locations = face_recognition.face_locations(image)

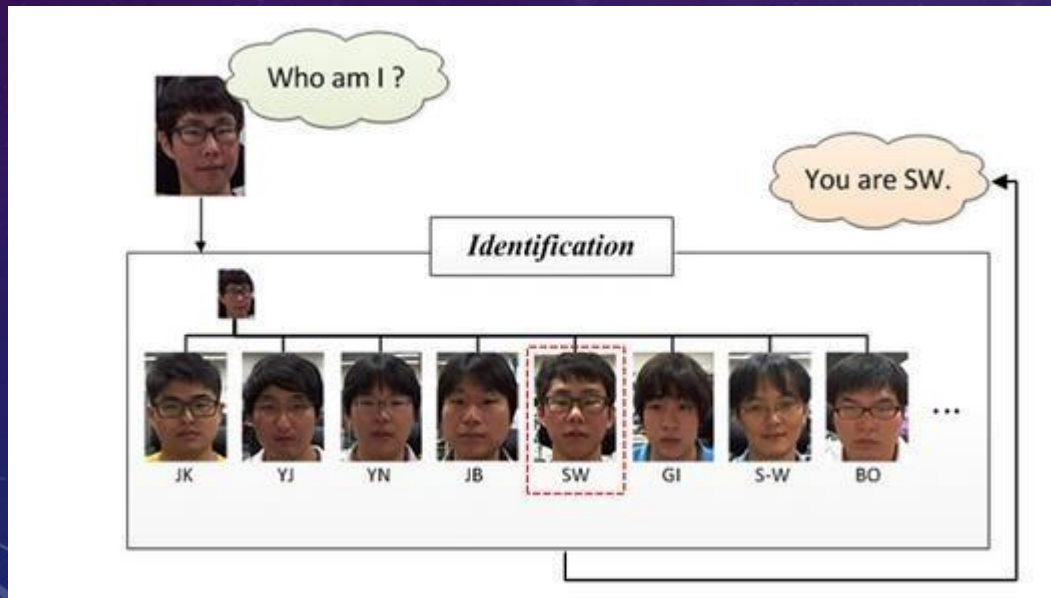
for face_location in face_locations:
    top, right, bottom, left = face_location

    face_image = image[top:bottom, left:right]
    pil_image = Image.fromarray(face_image)
    display(pil_image)
    pil_image.save(f'{top}.jpg')
```



Sol 7.1 : Face Recognition (3/10) – Face Identification

- Face identification is as shown in the figure below. It will answer "Who am I?"
- Compared with pair matching used in face verification, it uses classification more in the recognition stage.



```
# This is an example of running face recognition on a single image
# and drawing a box around each person that was identified.
import face_recognition
from PIL import Image, ImageDraw
import numpy as np

# Load a sample picture and get the feature.
obama_image = face_recognition.load_image_file("obama.jpg")
obama_face_encoding = face_recognition.face_encodings(obama_image)[0]

# Load a second sample picture and get the feature.
biden_image = face_recognition.load_image_file("biden.jpg")
biden_face_encoding = face_recognition.face_encodings(biden_image)[0]

# Create arrays of known face feature and their names
known_face_encodings = [
    obama_face_encoding,
    biden_face_encoding
]
known_face_names = [
    "Barack Obama",
    "Joe Biden"
]
```


Sol 7.1 : Face Recognition (4/10) – Face Identification

```
# Load an image with an unknown face
unknown_image = face_recognition.load_image_file("obama_group.jpg")

# Find all the faces and face features in the unknown image
face_locations = face_recognition.face_locations(unknown_image)
face_encodings = face_recognition.face_encodings(unknown_image, face_locations)

# Convert the image to a PIL-format image so that we can draw on top of it with the Pillow library
pil_image = Image.fromarray(unknown_image)
# Create a Pillow ImageDraw Draw instance to draw with
draw = ImageDraw.Draw(pil_image)

# Loop through each face found in the unknown image
for (top, right, bottom, left), face_encoding in zip(face_locations, face_encodings):
    # See if the face is a match for the known face(s)
    matches = face_recognition.compare_faces(known_face_encodings, face_encoding)

    name = "Unknown"

    # If a match was found in known_face_encodings, just use the first one.
    # if True in matches:
    #     first_match_index = matches.index(True)
    #     name = known_face_names[first_match_index]

    # Or instead, use the known face with the smallest distance to the new face
    face_distances = face_recognition.face_distance(known_face_encodings, face_encoding)
    best_match_index = np.argmin(face_distances)
    if matches[best_match_index]:
        name = known_face_names[best_match_index]
```

Sol 7.1 : Face Recognition (5/10) – Face Identification

```
# Draw a box around the face using the Pillow module
draw.rectangle(((left, top), (right, bottom)), outline=(0, 0, 255))

# Draw a label with a name below the face
text_width, text_height = draw.textsize(name)
draw.rectangle(((left, bottom - text_height - 10), (right, bottom)), fill=(0, 0, 255), outline=(0, 0, 255))
draw.text((left + 6, bottom - text_height - 5), name, fill=(255, 255, 255, 255))

# Remove the drawing library from memory as per the Pillow docs
del draw

# Display the resulting image
display(pil_image)

# You can also save a copy of the new image to disk if you want by uncommenting this line
pil_image.save("image_with_boxes.jpg")
```

Threshold:0.7



Sol 7.1 : Face Recognition (6/10) – Face Identification

Threshold:0.5



Threshold:0.35



Sol 7.1 : Face Recognition (7/10) – Function Explanation

```
def face_locations(img, number_of_times_to_upsample=1, model="hog"):
```

```
    """
```

Returns an array of bounding boxes of human faces in a image

:param img: An image (as a numpy array)

:param number_of_times_to_upsample: How many times to upsample the image looking for faces. Higher numbers find smaller faces.

:param model: Which face detection model to use. "hog" is less accurate but faster on CPUs. "cnn" is a more accurate deep-learning model which is GPU/CUDA accelerated (if available). The default is "hog".

:return: A list of tuples of found face locations in css (top, right, bottom, left) order

```
    """
```

```
    if model == "cnn":
```

```
        return [_trim_css_to_bounds(_rect_to_css(face.rect), img.shape) for face in _raw_face_locations(img,
number_of_times_to_upsample, "cnn")]
```

```
    else:
```

```
        return [_trim_css_to_bounds(_rect_to_css(face), img.shape) for face in _raw_face_locations(img,
number_of_times_to_upsample, model)]
```

Sol 7.1 : Face Recognition (8/10) – Function Explanation

```
def _raw_face_landmarks(face_image, face_locations=None, model="large"):
    if face_locations is None:
        face_locations = _raw_face_locations(face_image)
    else:
        face_locations = [_css_to_rect(face_location) for face_location in face_locations]

    pose_predictor = pose_predictor_68_point #dlib.shape_predictor(predictor_68_point_model)

    if model == "small":
        pose_predictor = pose_predictor_5_point #dlib.shape_predictor(predictor_5_point_model)

    return [pose_predictor(face_image, face_location) for face_location in face_locations]
```


Sol 7.1 : Face Recognition (9/10) – Function Explanation

```
def face_encodings(face_image, known_face_locations=None, num_jitters=1, model="small"):
```

```
    """
```

Given an image, return the 128-dimension face encoding for each face in the image.

:param face_image: The image that contains one or more faces

:param known_face_locations: Optional - the bounding boxes of each face if you already know them.

:param num_jitters: How many times to re-sample the face when calculating encoding. Higher is more accurate, but slower (i.e. 100 is 100x slower)

:param model: Optional - which model to use. "large" (default) or "small" which only returns 5 points but is faster.

:return: A list of 128-dimensional face encodings (one for each face in the image)

```
    """
```

```
    raw_landmarks = _raw_face_landmarks(face_image, known_face_locations, model)
```

```
    return [np.array(face_encoder.compute_face_descriptor(face_image, raw_landmark_set, num_jitters)) for  
            raw_landmark_set in raw_landmarks]
```

```
    #face_encoder = dlib.face_recognition_model_v1(face_recognition_model)
```

Sol 7.1 : Face Recognition (10/10) – Function Explanation

```
def face_distance(face_encodings, face_to_compare):
```

```
    """
```

Given a list of face encodings, compare them to a known face encoding and get a **euclidean** distance for each comparison face. The distance tells you how similar the faces are.

:param faces: List of face encodings to compare

:param face_to_compare: A face encoding to compare against

:return: A numpy ndarray with the distance for each face in the same order as the 'faces' array

```
    """
```

```
    if len(face_encodings) == 0:
```

```
        return np.empty((0))
```

```
    return np.linalg.norm(face_encodings - face_to_compare, axis=1)
```



Exercise 7.1 : Face Recognition

Please use `face_recognition.ipynb` from Moodle to solve the following problems

1. Face identification : Choose 10 subjects and one img/subject from the web to make a gallery set, and make a test set of 10 images which show different poses, lightings and expressions, and contain 2~5 subjects in the gallery set. Use different thresholds to plot the FRR v.s. FAR
2. Face verification : Compare your own photos and the images of celebrities on the web, use different thresholds for the verification test and show the corresponding feature distances & results.


Please upload your code and results to Moodle

FACE RECOGNITION + ATTENDANCE PROJECT



https://youtu.be/sz25xxF_AVE [52:23]

Machine Learning is Fun! Part 4: Modern Face Recognition with Deep Learning



One of these people is Will Farrell. The other is Chad Smith. I swear they are different people!

<https://medium.com/@ageitgey/machine-learning-is-fun-part-4-modern-face-recognition-with-deep-learning-c3cffc121d78>

Pixel-Level Recognition



Pixel-level Recognition

CVPR 2020 Tutorial

Alexander Kirillov
Facebook AI Research (FAIR)

FACEBOOK AI

<https://youtu.be/QCtHGT68RIU> [28:56]

LOSS FUNCTION IMPLEMENTATION



Example 7.2 : Loss function implementation

Prepare image for your task



Create Training / Validation / Testing List



Modify the model



Change different loss function



Test

Installation Process

- Download **loss_function.ipynb** from Moodle and open it
- Download **CASIA-DATA.zip** and unzip it to “./Database/”

```
%cd /content/DSS_Loss_Function
!unzip ./Database.zip
!wget https://www.dropbox.com/s/tdkptvd36nvx6ox/CASIA-Data.zip?dl=0 -O ./Database/CASIA-Data.zip
!unzip ./Database/CASIA-Data.zip -d ./Database
```



- Images are picked from CASIA FACE database.
 - 20 subjects
 - Label: 0~20

Database: CASIA-WebFace

- Unconstrained environment
- 10,575 identities with 494,414 images
- Collected from Internet Movie Database (IMDB)
- Large training data in public



Create Training/Validation/Testing List

Find create_csv.py

 create_csv.py

Use command to run create.py

```
!python create_csv.py
```

After execution, you can find “train.csv”, “test.csv” and “val.csv” under the “DataList” folder.

 DataList

-  test.csv
-  train.csv
-  val.csv

Create Training/Validation/Testing List

- Split training images in each subject for **Training, Validation and Test**.

```
import os
import csv
idx = 0
person = 0
with open(R'./DataList/train.csv','w', newline = '') as train:
    with open(R'./DataList/val.csv','w', newline = '') as val:
        with open(R'./DataList/test.csv','w', newline = '') as test:
            csv_writer = csv.writer(train)
            csv_writer.writerow(['image','id'])
            csv_writer1 = csv.writer(val)
            csv_writer1.writerow(['image','id'])
            csv_writer2 = csv.writer(test)
            csv_writer2.writerow(['image','id'])
            for root, dirs, filenames in os.walk(r'.\Database\CASIA-Data'):
                for k in dirs:
                    for root1, dirs1, filenames1 in os.walk(r'.\Database\CASIA-Data\{ }'.format(k)):
                        for i in filenames1:
                            idx += 1
                            data_path = os.path.join(root1,i)
                            print(idx)
                            if len(filenames1)*0 < idx <= len(filenames1)*0.8:
                                csv_writer.writerow([data_path,person])
                            if len(filenames1)*0.8 < idx <= len(filenames1)*0.9:
                                csv_writer1.writerow([data_path,person])
                            if len(filenames1)*0.9 < idx <= len(filenames1)*1.0:
                                csv_writer2.writerow([data_path,person])
            idx = 0
            person += 1
```

CASIA-DATA path

80% for Training
10% for Validation
10% for Test

Create Training/Validation/Testing List

This is “train.csv”.

The format is the same as “test.csv” and, “val.csv”

	A	B
1	image	id
2	.\Database\CASIA-Data\0000045\001.bmp	0
3	.\Database\CASIA-Data\0000045\002.bmp	0
4	.\Database\CASIA-Data\0000045\003.bmp	0
5	.\Database\CASIA-Data\0000045\004.bmp	0
6	.\Database\CASIA-Data\0000045\005.bmp	0
7	.\Database\CASIA-Data\0000045\006.bmp	0
8	.\Database\CASIA-Data\0000045\007.bmp	0
9	.\Database\CASIA-Data\0000045\008.bmp	0
10	.\Database\CASIA-Data\0000045\009.bmp	0
11	.\Database\CASIA-Data\0000045\011.bmp	0
12	.\Database\CASIA-Data\0000045\012.bmp	0
13	.\Database\CASIA-Data\0000099\001.bmp	1
14	.\Database\CASIA-Data\0000099\002.bmp	1
15	.\Database\CASIA-Data\0000099\003.bmp	1
16	.\Database\CASIA-Data\0000099\004.bmp	1
17	.\Database\CASIA-Data\0000099\005.bmp	1

Images path

Images label

Now, open the “Main_Angular.py” Let’s check the code !

Model Parameter

All the argument value can be found here, and the explanation will be shown in **help**

```
94 if __name__ == "__main__":
95
96     parser = argparse.ArgumentParser(description='DR_GAN')
97     # learning & saving parameters
98     parser.add_argument('-train', action='store_true', default=True,
99                         help='Generate pose modified image from given image')
100    parser.add_argument('-lr', type=float, default=0.0001, help='initial learning rate [default: 0.0002]')
101    parser.add_argument('-step-learning', action='store_true', default=False, help='enable lr step learning')
102    parser.add_argument('-lr-decay', type=float, default=0.1, help='initial decay learning rate [default: 0.1]')
103    parser.add_argument('-lr-step', type=int, default=35, help='Set Step to change lr by multiply lr-decay thru every lr-step epoch [default: 35]')
104    parser.add_argument('-beta1', type=float, default=0.5, help='adam optimizer parameter [default: 0.5]')
105    parser.add_argument('-beta2', type=float, default=0.999, help='adam optimizer parameter [default: 0.999]')
106    parser.add_argument('-epochs', type=int, default=5, help='number of epochs for train [default: 1000]')
107    parser.add_argument('-Train-Batch', type=int, default=8, help='batch size for training [default: 64]')
108    parser.add_argument('-Val-Batch', type=int, default=32, help='batch size for training [default: 4]')
109    parser.add_argument('-Test-Batch', type=int, default=32, help='batch size for training [default: 64]')
110    parser.add_argument('-snapshot-dir', type=str, default='snapshot', help='where to save the snapshot while training')
111    parser.add_argument('-save-freq', type=int, default=1, help='save learned model for every "-save-freq" epoch')
112    parser.add_argument('-cuda', action='store_true', default=False, help='enable the gpu')
113    parser.add_argument('-start-epoch', default=1, type=int, metavar='N', help='manual epoch number (useful on restarts)')
```


Model Parameter

```
114 # data source
115 parser.add_argument('-data-place', type=str, default=None, help='prepared data path to run program')
116 parser.add_argument('-output', type=str, default='Output', help='Output path for features')
117 parser.add_argument('-train-csv-file', type=str, default=None, help='csv file to load image for training')
118 parser.add_argument('-val-csv-file', type=str, default=None, help='csv file to load image for validation')
119 parser.add_argument('-test-csv-file', type=str, default=None, help='csv file to load image for test')
120 parser.add_argument('-Nd', type=int, default=10, help='initial Number of ID [default: 188]')
121 parser.add_argument('-Channel', type=int, default=3, help='initial Number of Channel [default: 3 (RGB Three Channel)]')
122 # option
123 parser.add_argument('-snapshot', type=str, default=None, help='filename of model snapshot(snapshot/(Single or Multiple)/(date)/(epoch)) [default: None]')
124 parser.add_argument('-test', action='store_true', default=None, help='Generate pose modified image from given image')
125 parser.add_argument('-resume', default='', type=str, metavar='PATH', help='path to latest checkpoint (default: none)')
126 parser.add_argument('-Angle-Loss', action='store_true', default=False, help='Use Angle Loss')
127 parser.add_argument('-pretrain', action='store_true', default=False)
128
129 args = parser.parse_args()
130 writer = SummaryWriter()
```

Whether use Angular Softmax Loss or not, the default setting is Softmax with cross entropy loss

Training Mode

Choose the training mode in argument value

```
149 elif args.train:
150     print("Parameters:")
151     for attr, value in sorted(args.__dict__.items()):
152         text = "\t{}={}\n".format(attr.upper(), value)
153         print(text)
154         with open('{}Parameters.txt'.format(args.snapshot_dir), 'a') as f:
155             f.write(text)
156
157     if args.train_csv_file is None or args.val_csv_file is None:
158         print(">>> Sorry, please set csv-file for your training/validation data")
159         exit()
160
161     else:
162         Model = VGG16(args)
163         print(Model)
164         Train(Model, args)
165
```

Set VGG16 as Model and print out the architecture

VGG-16

- Check DSS_Loss_Function/model/VGG16_Model.py

```
103 class VGG16(nn.Module):
104     def __init__(self, args, init_weights=True):
105         super(VGG16, self).__init__()
106         self.features = []
107         self.Nd = args.Nd
108         self.Channel = args.Channel
109         self.AngleLoss = args.Angle_Loss
110
111         ConvBlock1 = [
112             nn.Conv2d(self.Channel, 64, 3, 1, 1), # conv1_1
113             nn.BatchNorm2d(64),
114             nn.ELU(),
115             nn.Conv2d(64, 64, 3, 1, 1), # conv1_2
116             nn.BatchNorm2d(64),
117             nn.ELU(),
118             nn.MaxPool2d(2, stride=2), # pool1
119         ]
120         ConvBlock2 = [
121             nn.Conv2d(64, 128, 3, 1, 1), # conv2_1
122             nn.BatchNorm2d(128),
123             nn.ELU(),
124             nn.Conv2d(128, 128, 3, 1, 1), # conv2_2
125             nn.BatchNorm2d(128),
126             nn.ELU(),
127             nn.MaxPool2d(2, stride=2), # pool2
128         ]
```

```
129
130         ConvBlock3 = [
131             nn.Conv2d(128, 256, 3, 1, 1), # conv3_1
132             nn.BatchNorm2d(256),
133             nn.ELU(),
134             nn.Conv2d(256, 256, 3, 1, 1), # conv3_2
135             nn.BatchNorm2d(256),
136             nn.ELU(),
137             nn.Conv2d(256, 256, 3, 1, 1), # conv3_3
138             nn.BatchNorm2d(256),
139             nn.ELU(),
140             nn.MaxPool2d(2, stride=2), # pool3
141         ]
142         ConvBlock4 = [
143             nn.Conv2d(256, 512, 3, 1, 1), # conv4_1
144             nn.BatchNorm2d(512),
145             nn.ELU(),
146             nn.Conv2d(512, 512, 3, 1, 1), # conv4_2
147             nn.BatchNorm2d(512),
148             nn.ELU(),
149             nn.Conv2d(512, 512, 3, 1, 1), # conv4_3
150             nn.BatchNorm2d(512),
151             nn.ELU(),
152             nn.MaxPool2d(2, stride=2), # pool4
153         ]
154         ConvBlock5 = [
155             nn.Conv2d(512, 512, 3, 1, 1), # conv5_1
156             nn.BatchNorm2d(512),
157             nn.ELU(),
158             nn.Conv2d(512, 512, 3, 1, 1), # conv5_2
159             nn.BatchNorm2d(512),
160             nn.ELU(),
161             nn.Conv2d(512, 512, 3, 1, 1), # conv5_3
```

VGG-16

- Check DSS_Loss_Function/model/VGG16_Model.py

Decide how does the data flow in structure

```
self.convLayers1 = nn.Sequential(*ConvBlock1)
self.convLayers2 = nn.Sequential(*ConvBlock2)
self.convLayers3 = nn.Sequential(*ConvBlock3)
self.convLayers4 = nn.Sequential(*ConvBlock4)
self.convLayers5 = nn.Sequential(*ConvBlock5)
```

```
self.FC6 = nn.Sequential(
    nn.Linear(2048, 4096),
    nn.ReLU(),
    nn.Dropout(0.5),
)
```

```
self.FC7 = nn.Sequential(
    nn.Linear(4096, 4096),
    nn.ReLU(),
    nn.Dropout(0.6),
)
```

```
if self.AngleLoss:
    self.FC8 = AngleLinear(4096, self.Nd)
else:
    self.FC8 = nn.Linear(4096, self.Nd)
```

```
194 def forward(self, input, ExtractMode=False):
195
196     x1 = self.convLayers1(input)
197     x2 = self.convLayers2(x1)
198     x3 = self.convLayers3(x2)
199     x4 = self.convLayers4(x3)
200     x = self.convLayers5(x4)
201
202     x = x.view(np.shape(x)[0], -1) # np.shape(x)[0] -> batch
203     x = self.FC6(x)
204     x = self.FC7(x)
205     self.features = x
206     x = self.FC8(x)
207
208     if ExtractMode:
209         return self.features
210     else:
211         return x
```

Feature extraction

Introduction VGG-16

Double
ConvBlock

```
VGG16(  
  (convLayers1): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ELU(alpha=1.0)  
    (3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): ELU(alpha=1.0)  
    (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (convLayers2): Sequential(  
    (0): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ELU(alpha=1.0)  
    (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): ELU(alpha=1.0)  
    (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (convLayers3): Sequential(  
    (0): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ELU(alpha=1.0)  
    (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): ELU(alpha=1.0)  
    (6): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (8): ELU(alpha=1.0)  
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (convLayers4): Sequential(  
    (0): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ELU(alpha=1.0)  
    (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): ELU(alpha=1.0)  
    (6): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (8): ELU(alpha=1.0)  
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (convLayers5): Sequential(  
    (0): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ELU(alpha=1.0)  
    (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): ELU(alpha=1.0)  
    (6): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (8): ELU(alpha=1.0)  
    (9): AdaptiveAvgPool2d(output_size=(2, 2))  
  )  
  (FC6): Sequential(  
    (0): Linear(in_features=2048, out_features=4096, bias=True)  
    (1): ReLU()  
    (2): Dropout(p=0.5)  
  )  
  (FC7): Sequential(  
    (0): Linear(in_features=4096, out_features=4096, bias=True)  
    (1): ReLU()  
    (2): Dropout(p=0.6)  
  )  
  (FC8): Linear(in_features=4096, out_features=10, bias=True)  
)
```

Double
ConvBlock

Triple
ConvBlock

Triple
ConvBlock

Triple
ConvBlock

Triple
FC Layer

Training Process

```
28 def Train(Model, args):
29
30     Nd = args.Nd
31     beta1_Adam = args.beta1
32     beta2_Adam = args.beta2
33
34     if args.cuda:
35         Model.cuda()
36
37     optimizer = optim.Adam(Model.parameters(), lr=args.lr, betas=(beta1_Adam, beta2_Adam))
38     Model.train()
39     steps = 0
40     CUDNN.benchmark = True
41
42     for epoch in range(args.start_epoch, args.epochs+1):
43
44         if args.step_learning:
45             adjust_learning_rate(optimizer, epoch, args)
46
47         transformed_dataset = FacIdPoseDataset(args.train_csv_file, transform=transforms.Compose(
48             [transforms.Resize(32),
49              transforms.RandomCrop(28),
50              transforms.ToTensor()])))
51         dataloader = DataLoader(transformed_dataset, batch_size=args.Train_Batch, shuffle=True)
```

Decide the way to optimize model (default : Adam)

Original CASIA input is 112x112, resize into 256x256 and random crop back to 224x224 as data augmentation.

Training Process

```
53 for i, batch_data in enumerate(dataloader):
54     Model.zero_grad()
55     batch_image = torch.FloatTensor(batch_data[0].float())
56     batch_id_label = batch_data[2]
57     if args.cuda:
58         batch_image, batch_id_label = batch_image.cuda(), batch_id_label.cuda()
59     batch_image, batch_id_label = Variable(batch_image), Variable(batch_id_label)
60
61     steps += 1
62
63     Prediction = Model(batch_image)
64     Loss = Model.ID_Loss(Prediction, batch_id_label)
65
66     Loss.backward()
67     optimizer.step()
68     log_learning(epoch, steps, 'VGG16_Model', args.lr, Loss.data, args)
69     writer.add_scalar('Train/Train_Loss', Loss, steps)
70     # Validation_Process(Model, epoch, writer, args)
71     Validation_Process(Model, epoch, writer, args)
```

Send the image to Model, and output the prediction.

Loss function

Define the loss function, we already prepare two kinds of loss function for you guys

```
100 loss_criterion = nn.CrossEntropyLoss().cuda()  
101 loss_criterion_Angular = AngleLoss().cuda()
```

Given the FC prediction and label, than we can calculate the loss

```
213 def ID_Loss(self, predic, label):  
214  
215     if self.AngleLoss:  
216         Loss = loss_criterion_Angular((predic[0][:, :self.Nd], predic[1][:, :self.Nd]), label)  
217     else:  
218         Loss = loss_criterion(predic[:, :self.Nd], label)  
219  
220  
221     return Loss
```

Train the Model

Direct to the file path and input the command in the Terminal :

(Please build pytorch-gpu with the following command

```
conda install pytorch==1.0.0 torchvision==0.2.1 cuda80 -c pytorch)
```

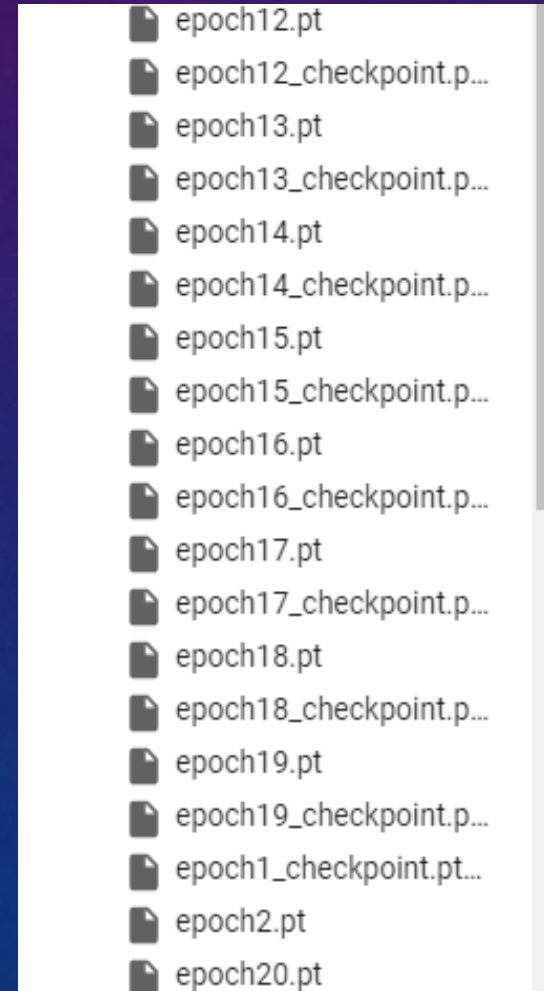

```
!python Main_Angular.py -cuda -train -data-place=./Database/CASIA-Data -train-csv-  
file=/content/DSS_Loss_Function/DataList/train.csv -val-csv-  
file=/content/DSS_Loss_Function/DataList/val.csv -Nd=20 -Train-Batch=8
```

```
Fri, 22 Nov 2019 11:46:25 +0000 EPOCH : 1, step : 259, VGG16_Model_Lr: 0.0001, VGG16_Model, 3.203601598739624  
Fri, 22 Nov 2019 11:46:25 +0000 EPOCH : 1, step : 260, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.9001379013061523  
Fri, 22 Nov 2019 11:46:25 +0000 EPOCH : 1, step : 261, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.620979070663452  
Fri, 22 Nov 2019 11:46:25 +0000 EPOCH : 1, step : 262, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.2005977630615234  
Fri, 22 Nov 2019 11:46:25 +0000 EPOCH : 1, step : 263, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.6367361545562744  
Fri, 22 Nov 2019 11:46:25 +0000 EPOCH : 1, step : 264, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.5872464179992676  
Fri, 22 Nov 2019 11:46:25 +0000 EPOCH : 1, step : 265, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.5509464740753174  
Fri, 22 Nov 2019 11:46:25 +0000 EPOCH : 1, step : 266, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.907942533493042  
Fri, 22 Nov 2019 11:46:26 +0000 EPOCH : 1, step : 267, VGG16_Model_Lr: 0.0001, VGG16_Model, 3.0043704509735107  
Fri, 22 Nov 2019 11:46:26 +0000 EPOCH : 1, step : 268, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.466834545135498  
Fri, 22 Nov 2019 11:46:26 +0000 EPOCH : 1, step : 269, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.3059823513031006  
Fri, 22 Nov 2019 11:46:26 +0000 EPOCH : 1, step : 270, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.745974063873291  
Fri, 22 Nov 2019 11:46:26 +0000 EPOCH : 1, step : 271, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.6620829105377197  
Fri, 22 Nov 2019 11:46:26 +0000 EPOCH : 1, step : 272, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.74006986618042  
Fri, 22 Nov 2019 11:46:26 +0000 EPOCH : 1, step : 273, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.995265483856201  
Fri, 22 Nov 2019 11:46:26 +0000 EPOCH : 1, step : 274, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.3112964630126953  
Fri, 22 Nov 2019 11:46:27 +0000 EPOCH : 1, step : 275, VGG16_Model_Lr: 0.0001, VGG16_Model, 2.737488031387329  
Fri, 22 Nov 2019 11:46:28 +0000 EPOCH : 1, step : 276, VGG16_Model_Lr: 0.0001, VGG16_Model, 1.990052342414856  
Start Validating...  
>>> epoch: '1'  
>>> ID_Precision: '0.1423611111111111'  
>>> Validation_Loss: '2.692030191421509'
```

Training process look like

Model saved

The model will save every epoch, and save in VGG16/snapshot.



You will need this model in Test phase

Test Phase

Type the command below to extract the feature from the model

```
!python Main_Angular.py -cuda -test -Nd=20 -test-csv-file=/content/DSS_Loss_Function/DataList/test.csv -  
snapshot=./snapshot/2021-01-04_09-33-41/epoch20 -output=./Output
```

Change to your model path

```
>>> Loading model from [./snapshot/2019-11-22_11-45-44/epoch8]...  
Finish Processing 32 images...  
Finish Processing 64 images...  
Finish Processing 96 images...  
Finish Processing 128 images...  
Finish Processing 160 images...  
Finish Processing 192 images...  
Finish Processing 224 images...  
Finish Processing 256 images...  
Finish Processing 286 images...
```

Test Phase

Open the “Feature_Compare.py” in VGG16

```
import os
import sys
import numpy as np
import scipy.io as sio
from scipy.spatial.distance import cdist
from numpy.linalg import norm
import csv
def read_feature(path):
    elements1 = []
    with open(path) as file:
        for line in file:
            line = line.strip().split()
            elements1.append(line)
    feature = np.array(elements1)
    feature = feature[:,0].astype(np.float64)
    return feature
```

```
Feature_path = './Output/snapshot/2019-11-22_11-45-44/epoch8/Feature'
```

```
Feature_dir = os.listdir(Feature_path)
```

```
Feature_dir.sort(key=lambda x:int(x))
```

Change to your own feature path

In Test phase, we extract the feature vector,
(not the FC output layer)
and compare the distance between each sample.

We choose the first sample in each class as the gallery image, and compare with all probe (test) sample to check the minimize distance.

Test Phase

Open the “Feature_Compare.py” in VGG16

```
## Choose first sample as gallery
Gallery=[]
for ii in range(len(Feature_dir)):
    fea_file = os.listdir(os.path.join(Feature_path,Feature_dir[ii]))
    fea = read_feature(os.path.join(Feature_path,Feature_dir[ii],fea_file[0]))
    Gallery.append(fea.T)
Gallery = np.array(Gallery)
Acc = 0
counter = 0

## Compare distance
for ii in range(len(Feature_dir)):
    fea_file = os.listdir(os.path.join(Feature_path, Feature_dir[ii]))
    for jj in range(len(fea_file)):
        Probe = read_feature(os.path.join(Feature_path,Feature_dir[ii],fea_file[jj]))
        distance = cdist(Gallery, Probe.reshape(-1,1).T , 'euclidean')
        value = distance.min()
        position = np.where(distance == value)
        if position[0] == (ii):
            Acc = Acc+1
            counter = counter +1
accuracy = Acc/counter
print("The accuracy =
{ }%\n'.format(accuracy*100))
```

Pairwise distance between two sets of observation

Reference Performance

```
(C:\Users\Micky\Anaconda3\envs\pytorch-cpu) F:\DSS_2019\VGG_NEW\VGG16>python Feature_Compare.py
The accuracy = 28.321678321678323%
```


Exercise 7.2 : Loss function implementation

Please change the loss function to “Angular Softmax Loss” with the same setting you have done in “Sample 7.2”, and use feature comparison to calculate the performance.

1. Compare the performance between difference loss function (Angular Softmax/ Cross entropy with softmax function)
2. Try to convert the pairwise distance from “Euclidean” to “Cosine”

Write down your observation and upload the report in MS Word or PDF to Moodle

Exercise 7.3 : Age classification

Please change the loss function to “Angular Softmax Loss” with the same setting you have done in **Example 7.2**, and use feature comparison to calculate the performance.

1. Use an age database “age.zip” from Colab that divides age into four intervals for training an age classifier and split the dataset with 70% for training, 10% for validation, 20% for testing.
2. Compare the performance between difference loss functions (Angular Softmax/ Cross Entropy with Softmax function)
3. Try to convert the pairwise distance from “Euclidean” to “Cosine”

Hint : The parameter “Nd” depends on the number of categories you have

Elaborate your solution and upload it to Moodle. You are encouraged to use the data and define your own version of additional experiments, and merge them to the report.

Exercise 7.4 : Object classification

Please change the loss function to “Angular Softmax Loss” with the same setting you have done in **Example 7.2**, and use feature comparison to calculate the performance.

1. Use an object database “cifar10.zip” given from Colab that divides object into ten intervals for training an object classifier and split the dataset with 70% for training, 10% for validation, 20% for testing.
2. Compare the performance between difference loss functions (Angular Softmax/ Cross Entropy with Softmax function)
3. Try to convert the pairwise distance from “Euclidean” to “Cosine”

Hint : The parameter “Nd” depends on the number of categories you have

Elaborate your solution and upload it to Moodle. You are encouraged to use the data and define your own version of additional experiments, and merge them to the report.